

Welcome to the World of Distributed Messaging Queue

- The consumer pull data directly to the broker that is **the leader for the partition.**
- Is Consumer use zookeeper?
 - Yes (0.8.x)
 - No (0.9.x +)

- General messaging system has two model:
 - Queuing
 - Publish-Subscribe
- Queuing:
 - A pool of consumers will read from a server and each message goes to one of them.
- Publish-Subscribe:
 - The message is broadcast to all consumers.

Consumer

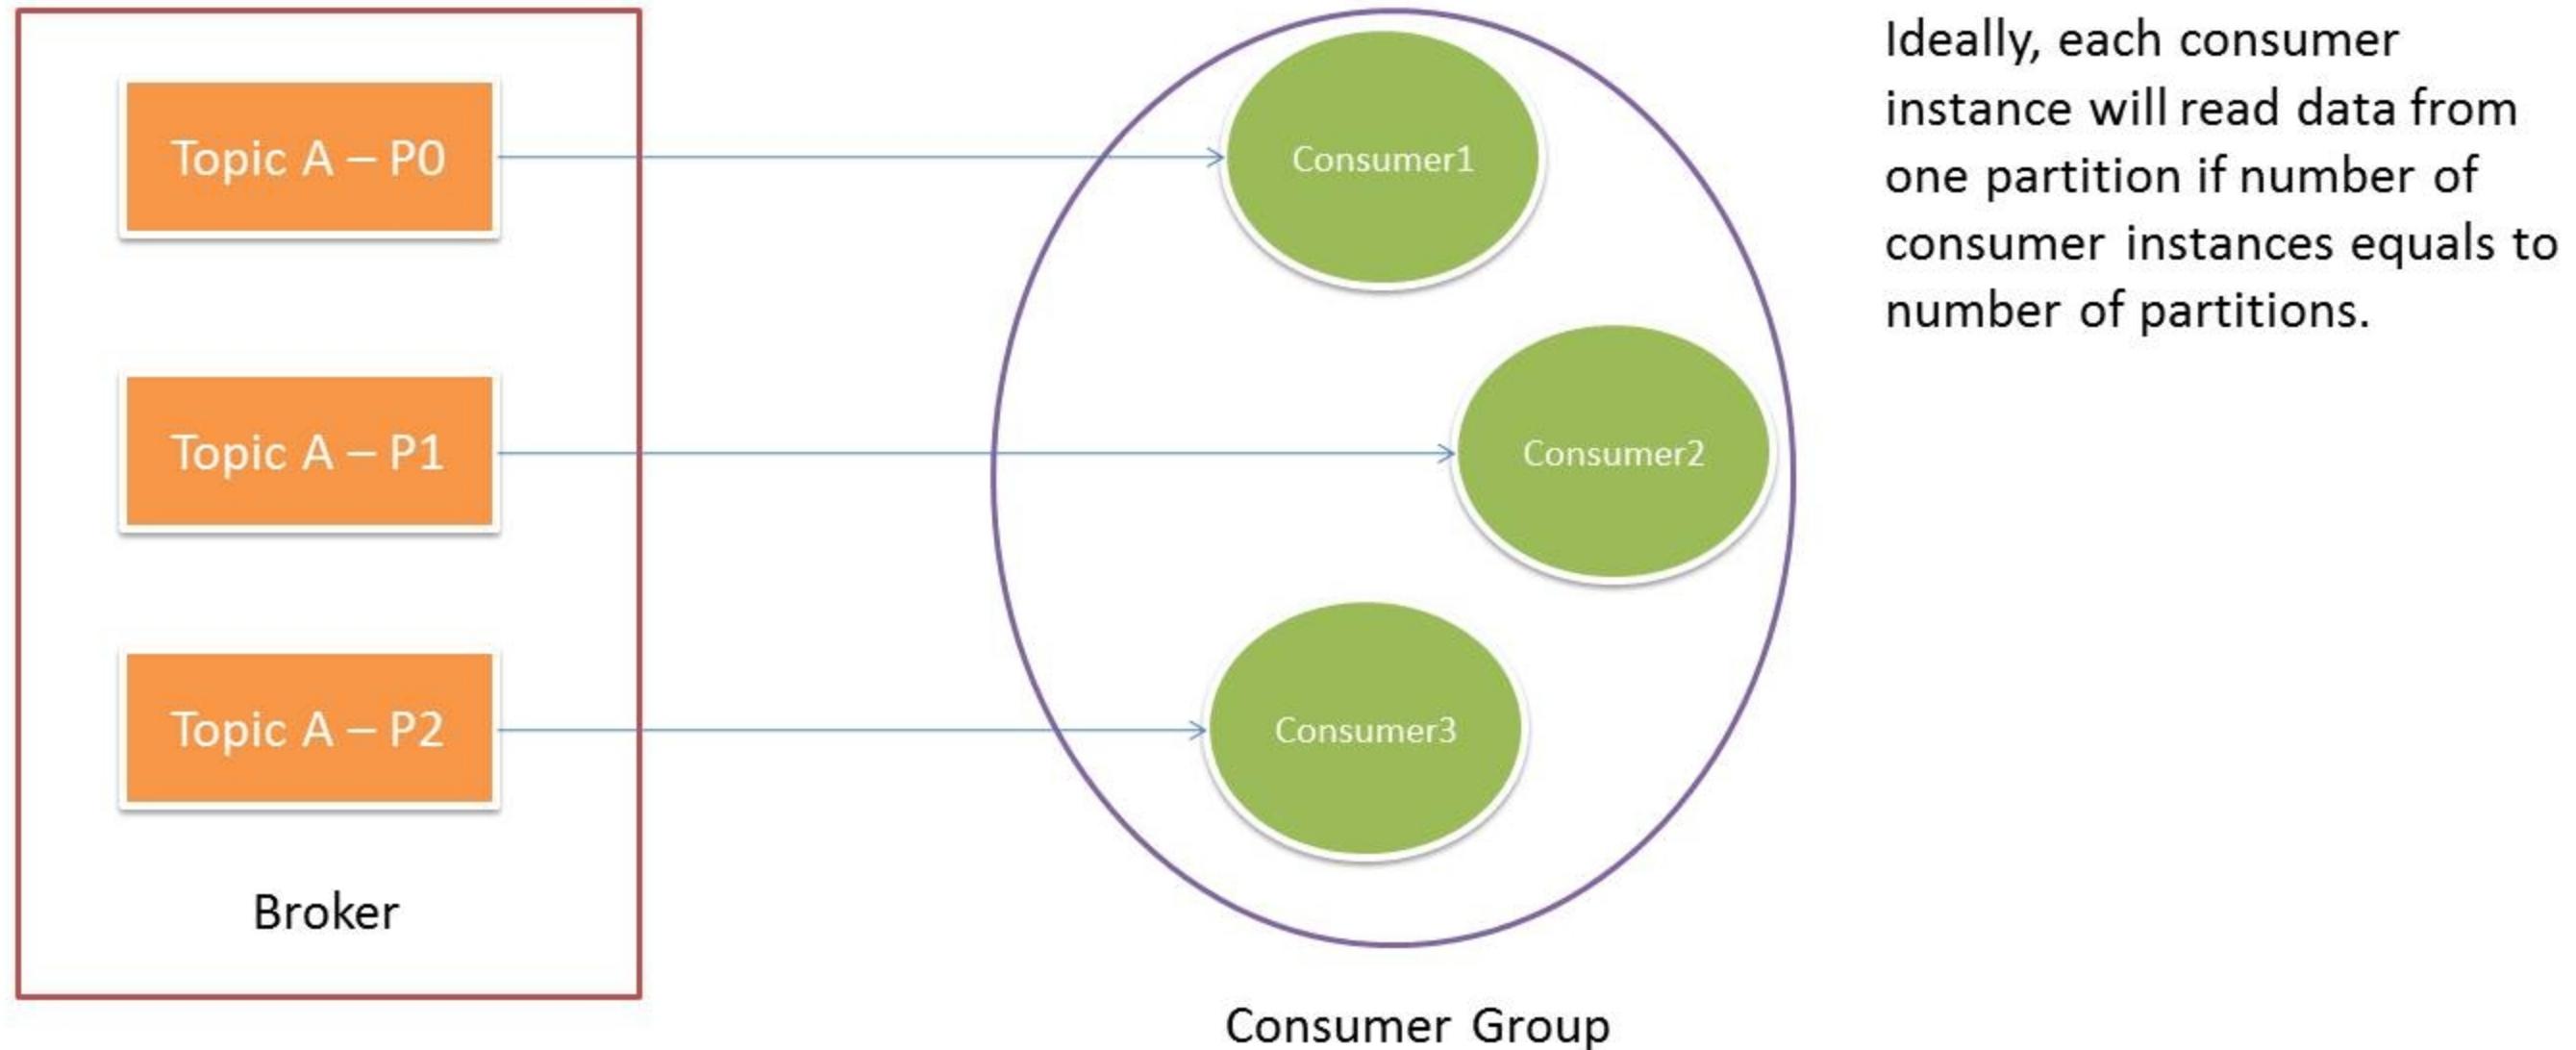


1. Which model Kafka consumer support?

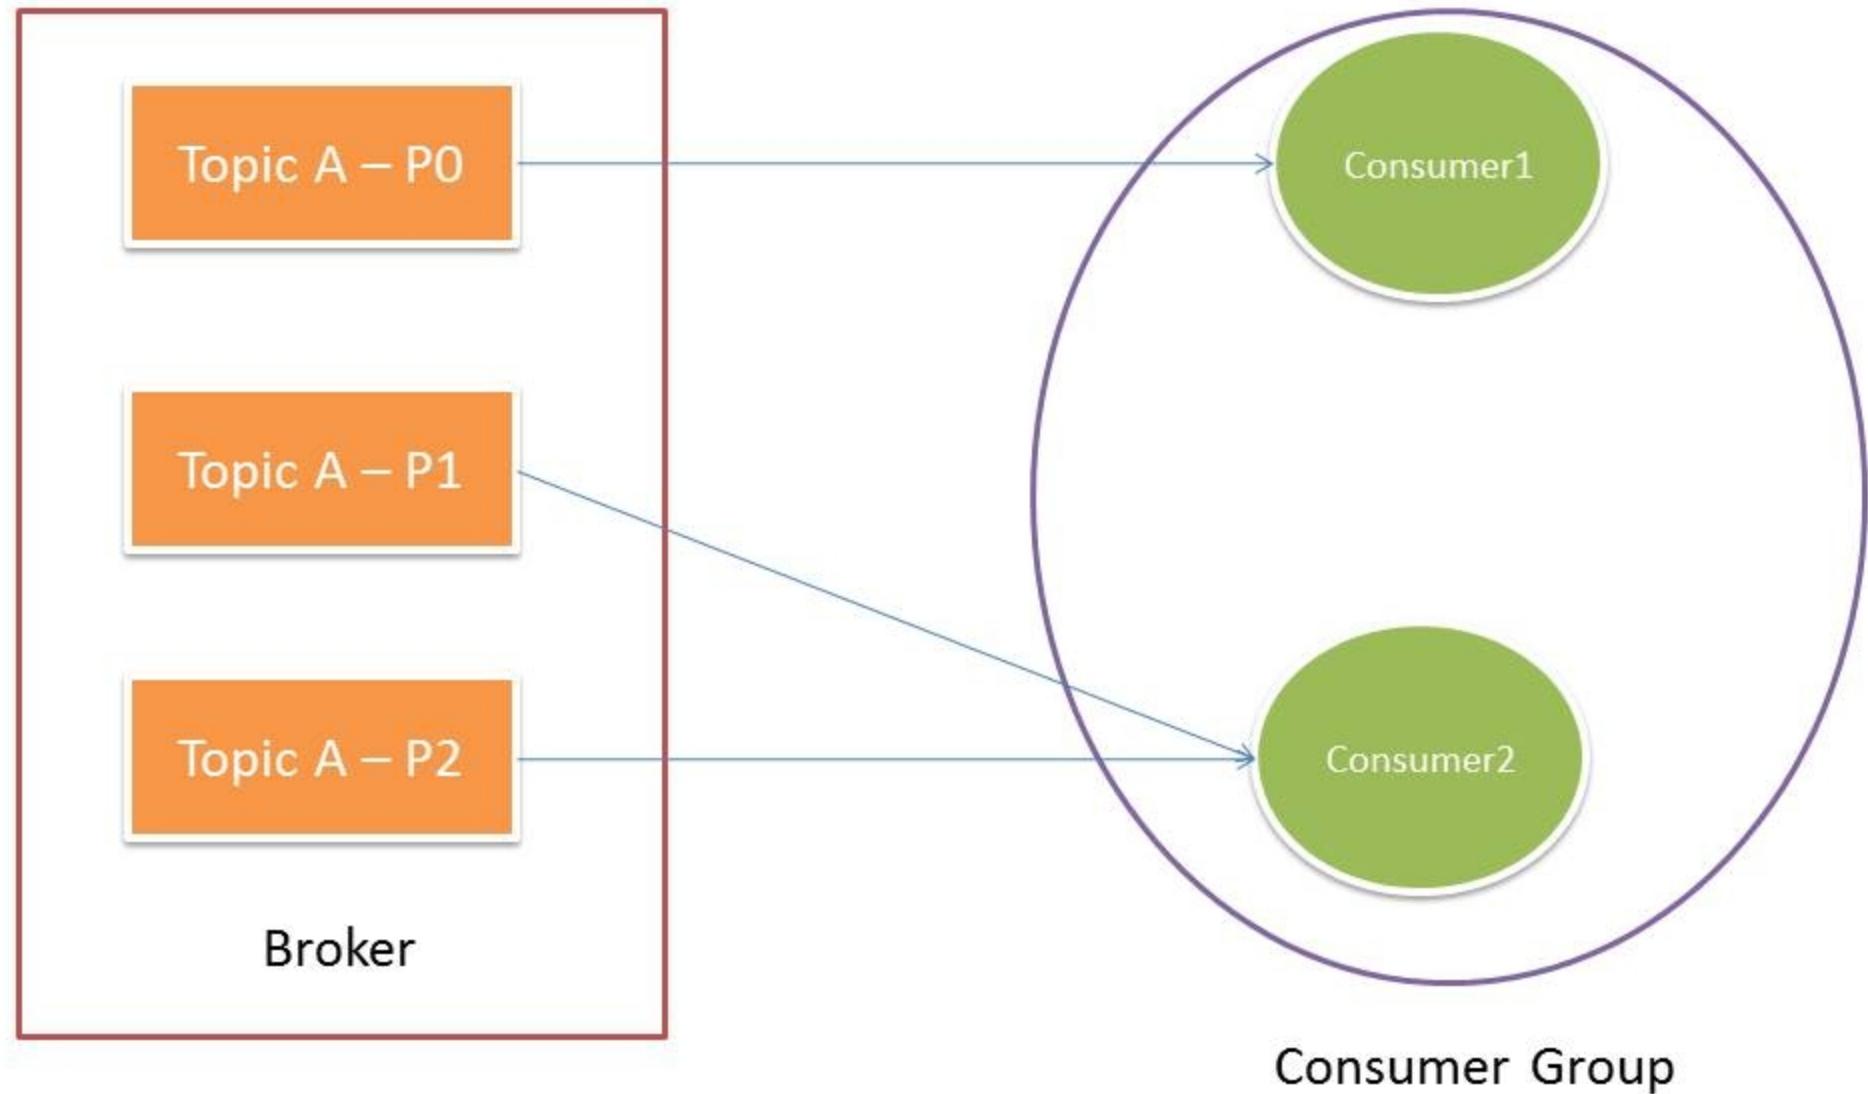
- Which model Kafka consumer support?
 - Kafka Consumer supports both queuing and publish-subscribe mode.
- How?
 - Each Kafka Consumer belongs to a consumer group
 - A one consumer group can have multiple consumer instance
 - Each message published to a topic is delivered to one consumer instance within each subscribing consumer group.
 - Consumer instances can be in separate processes or on separate machines.

- If all the consumer instances have the same consumer group, then this works just like a traditional queue balancing load over the consumers.
- If all the consumer instances have different consumer groups, then this works like publish-subscribe and all messages are broadcast to all consumers.

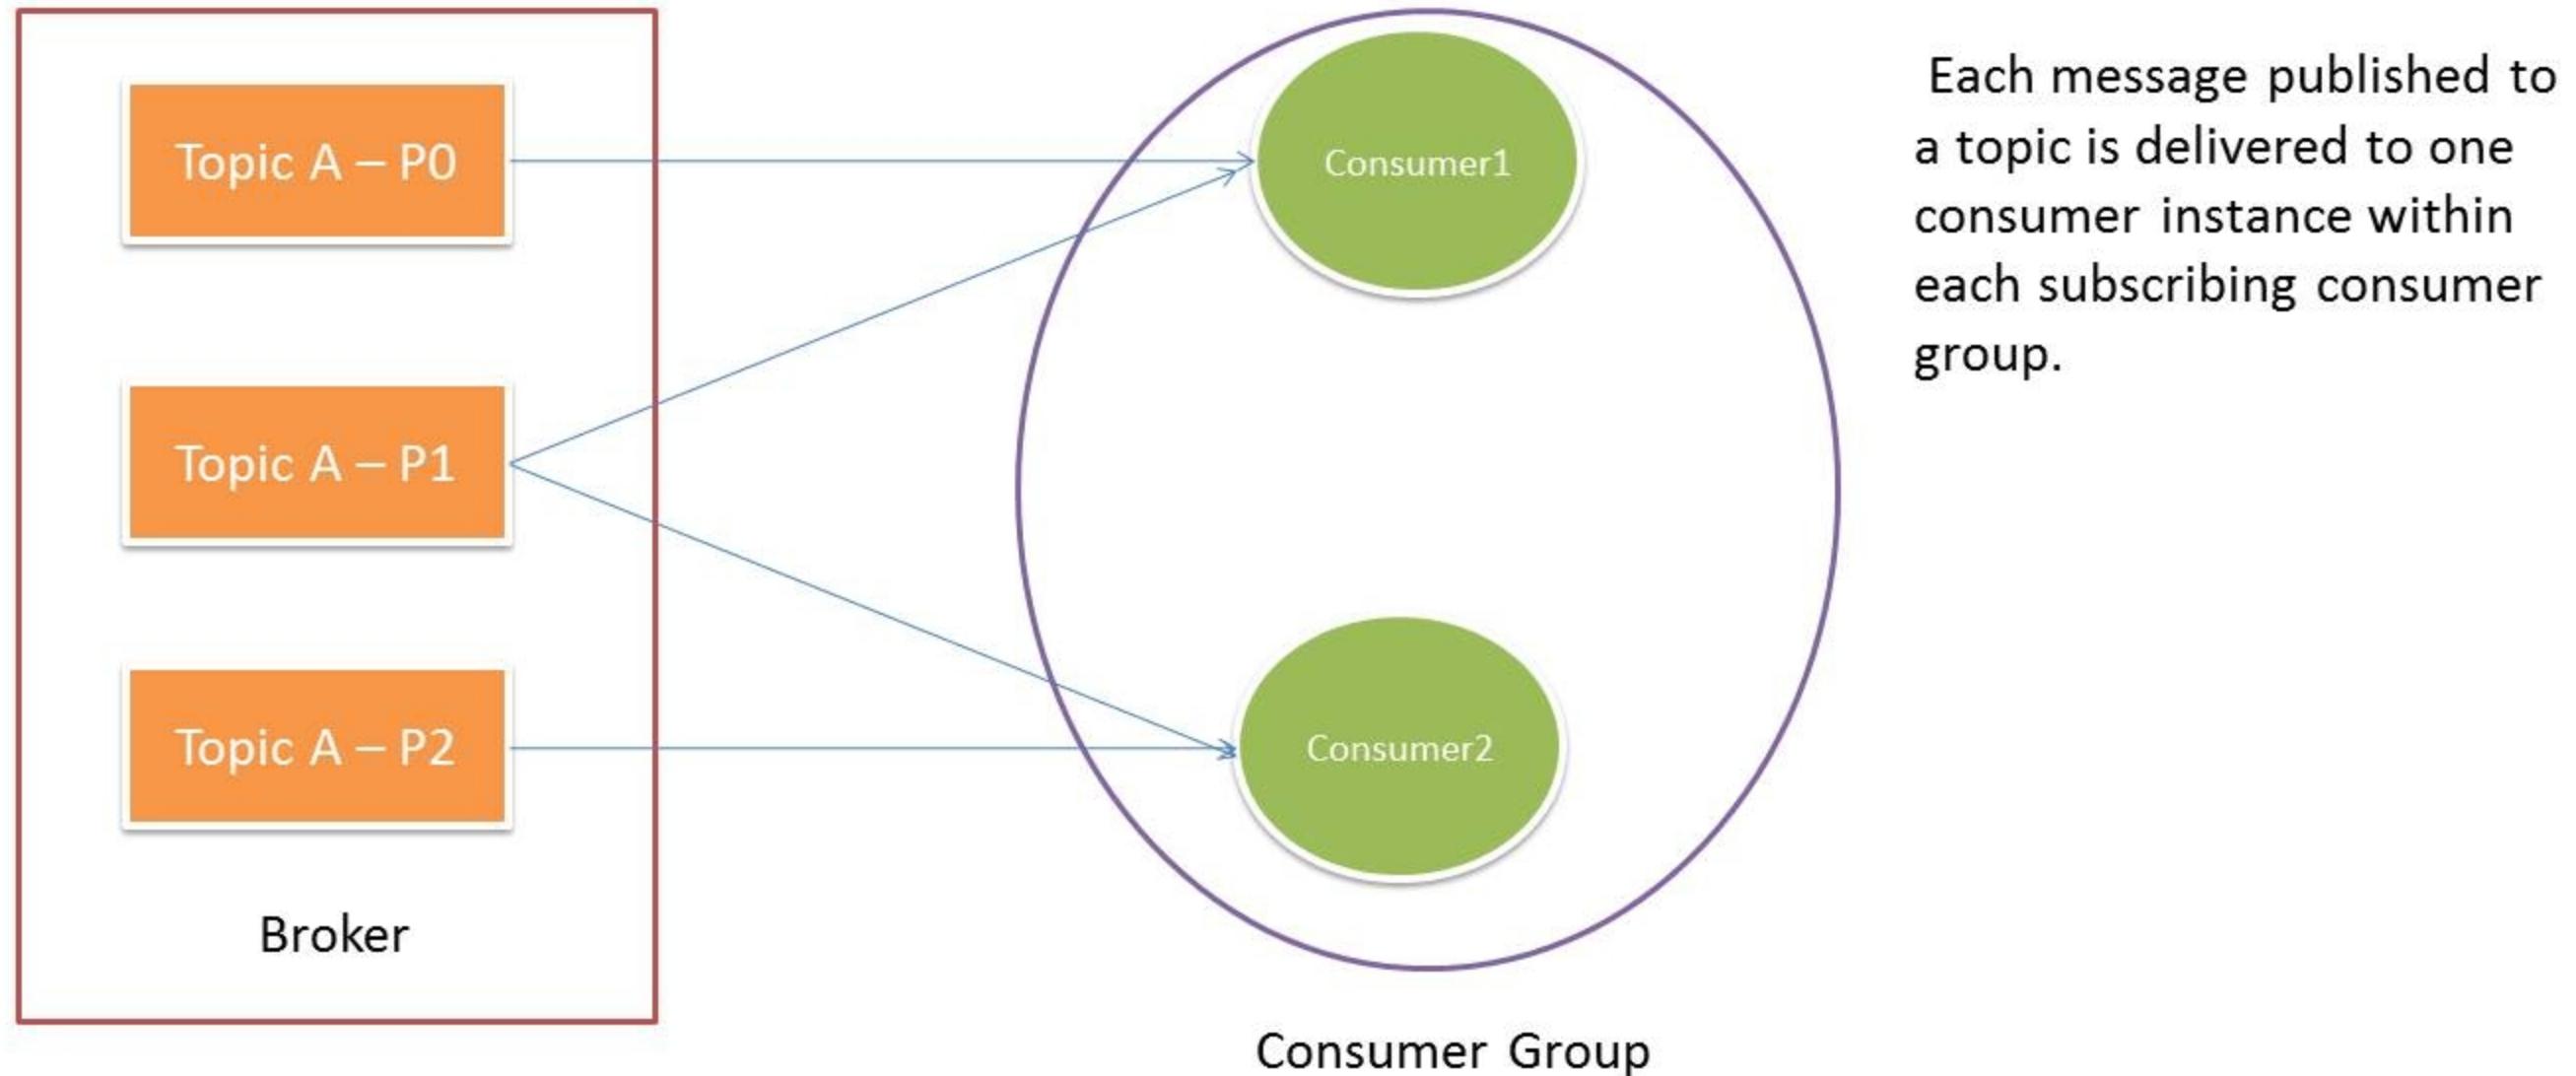
Consumer - Queuing



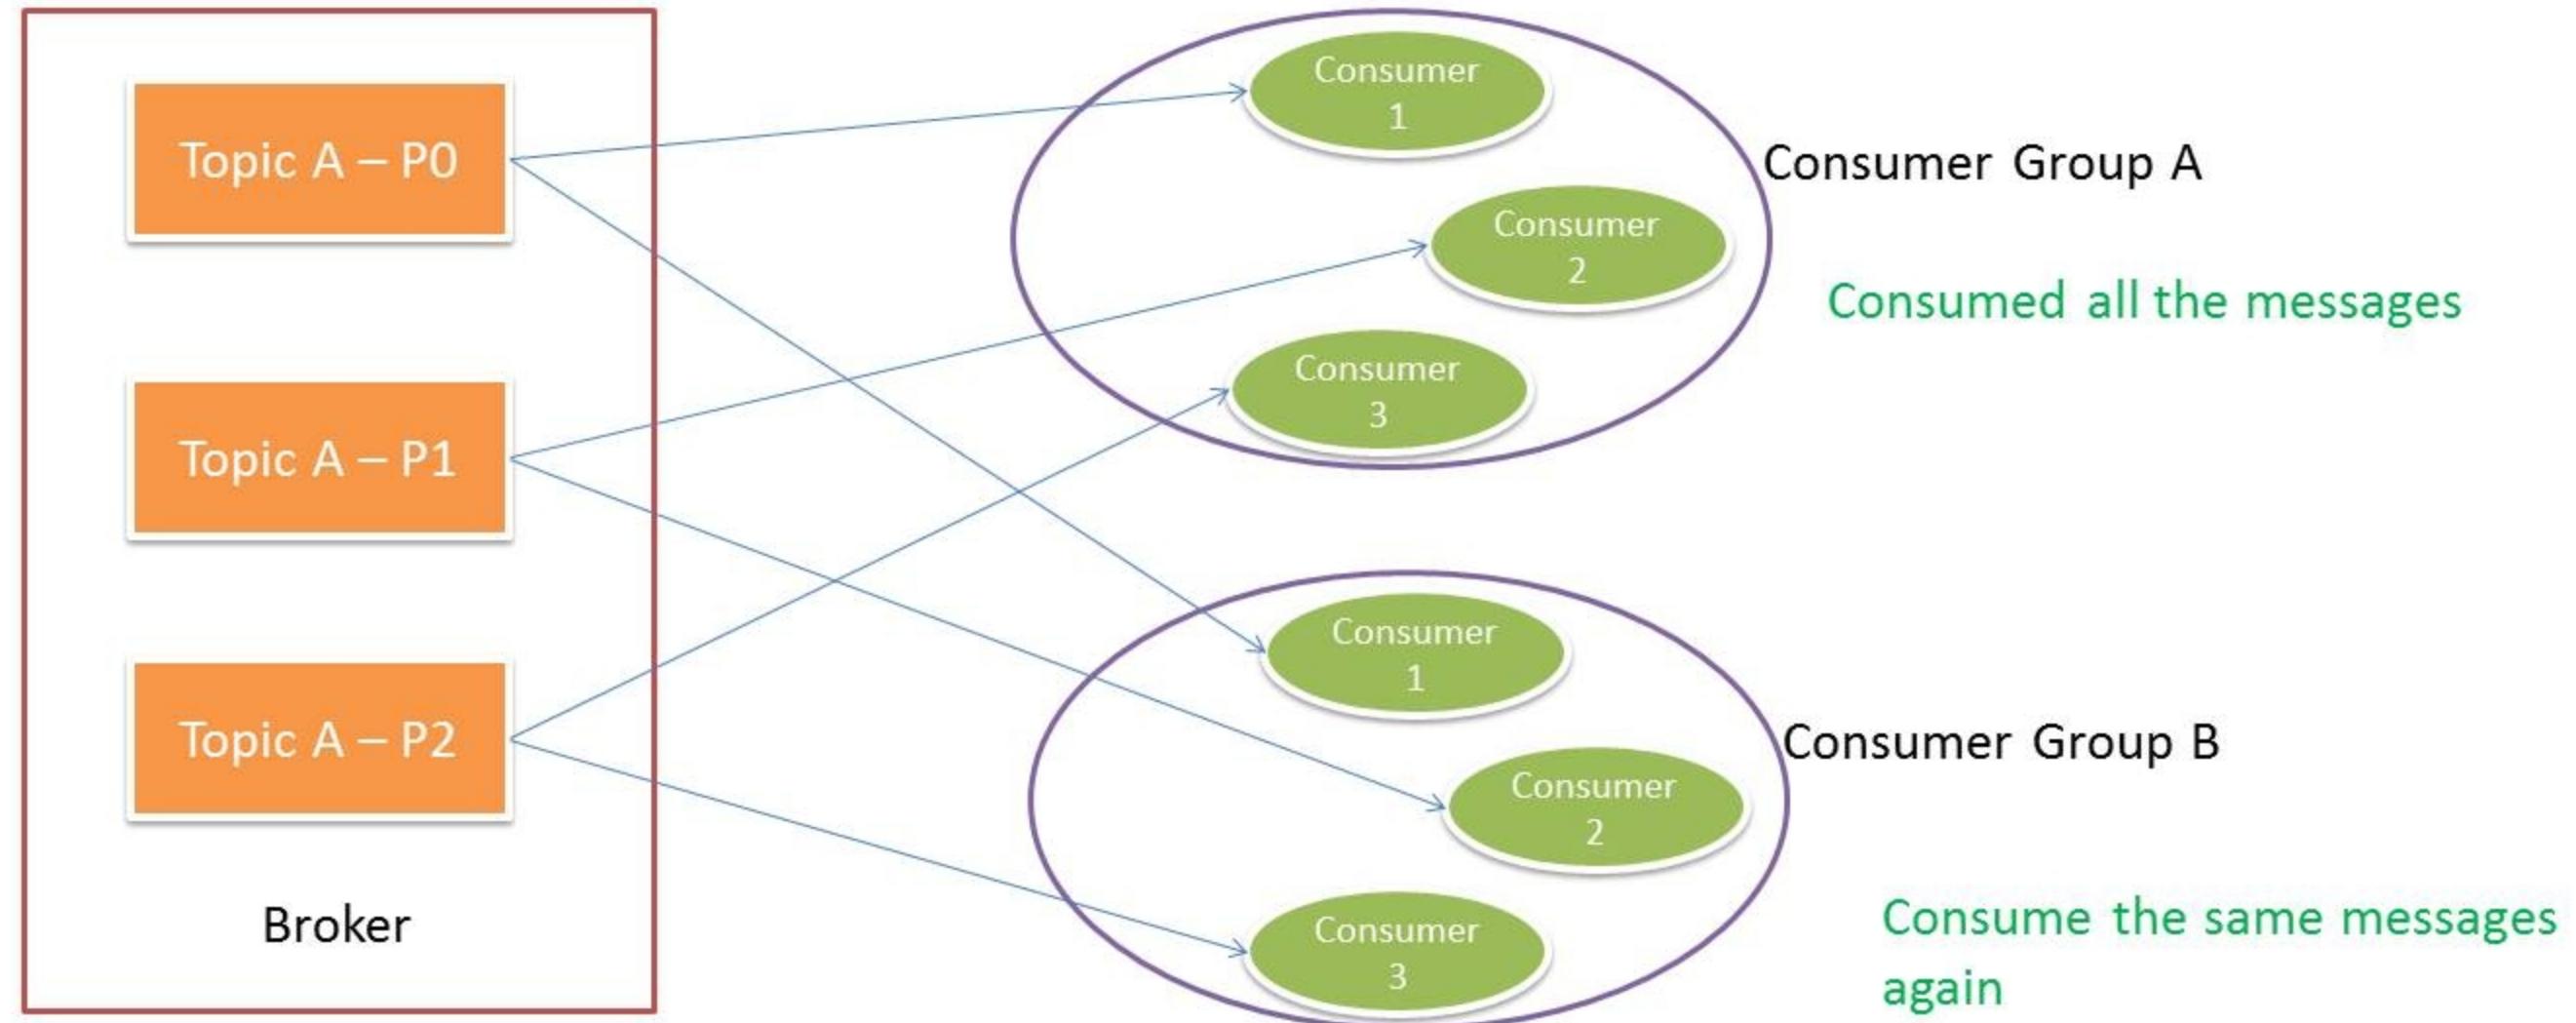
Consumer - Queuing



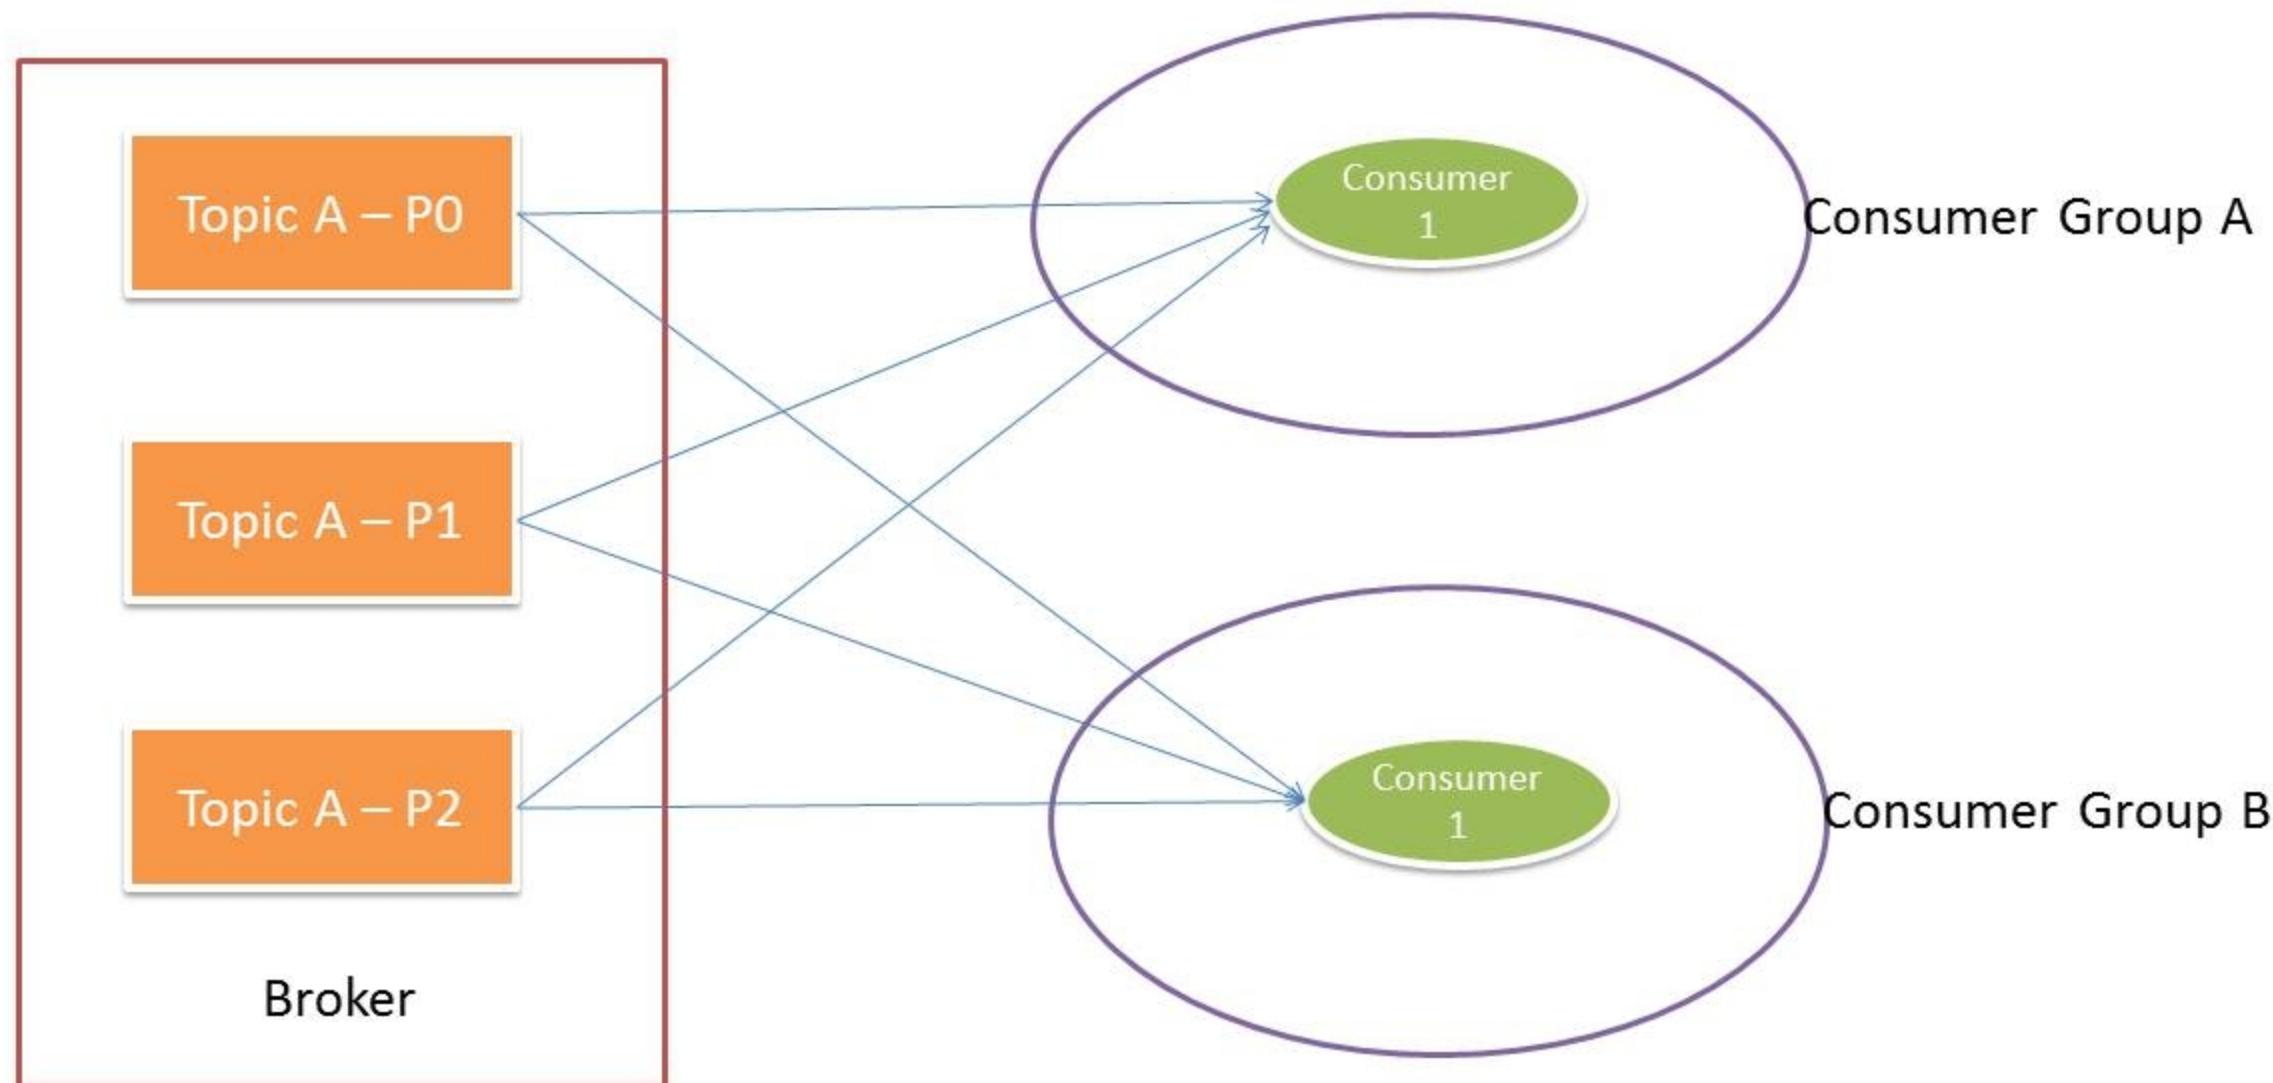
Consumer - Queuing



Consumer Group



Consumer Broadcast



- Push data into Kafka
 - bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test-replication-1
- Pull data from Kafka
 - bin/kafka-console-consumer.sh --zookeeper localhost:2181 --topic test-replication-1 --from-beginning

- Installation of kafka on multiple machines

- For 0.8.x
 - Producer Daemon
 - Python
 - Go (AKA golang)
 - C
 - C++
 - .net
 - Clojure
 - Ruby
 - Node.js
 - Alternative Java Clients
 - Storm
 - Scala DSL
 - HTTP REST
 - JRuby
 - Perl
 - stdin/stdout
 - Erlang

- Need to set the following properties:
`metadata.broker.list=broker1:9092,broker2:9092`
`serializer.class=kafka.serializer.StringEncoder`
`request.required.acks=1`
- The property “**metadata.broker.list**” is used to determine the leader partition for each topic.
- The property “**serializer.class**” defines what Serializer to use when preparing the message for transmission to the Broker.
- The property "**request.required.acks**" tells Kafka that you want your Producer to require an acknowledgement from the Broker that the message was received.

- API to create topic into Kafka.
 - Sample Example
- Produce data into Kafka
- How data distribute between partition

- Need to set the following properties:

zookeeper.connect=IP1:port,IP2:port

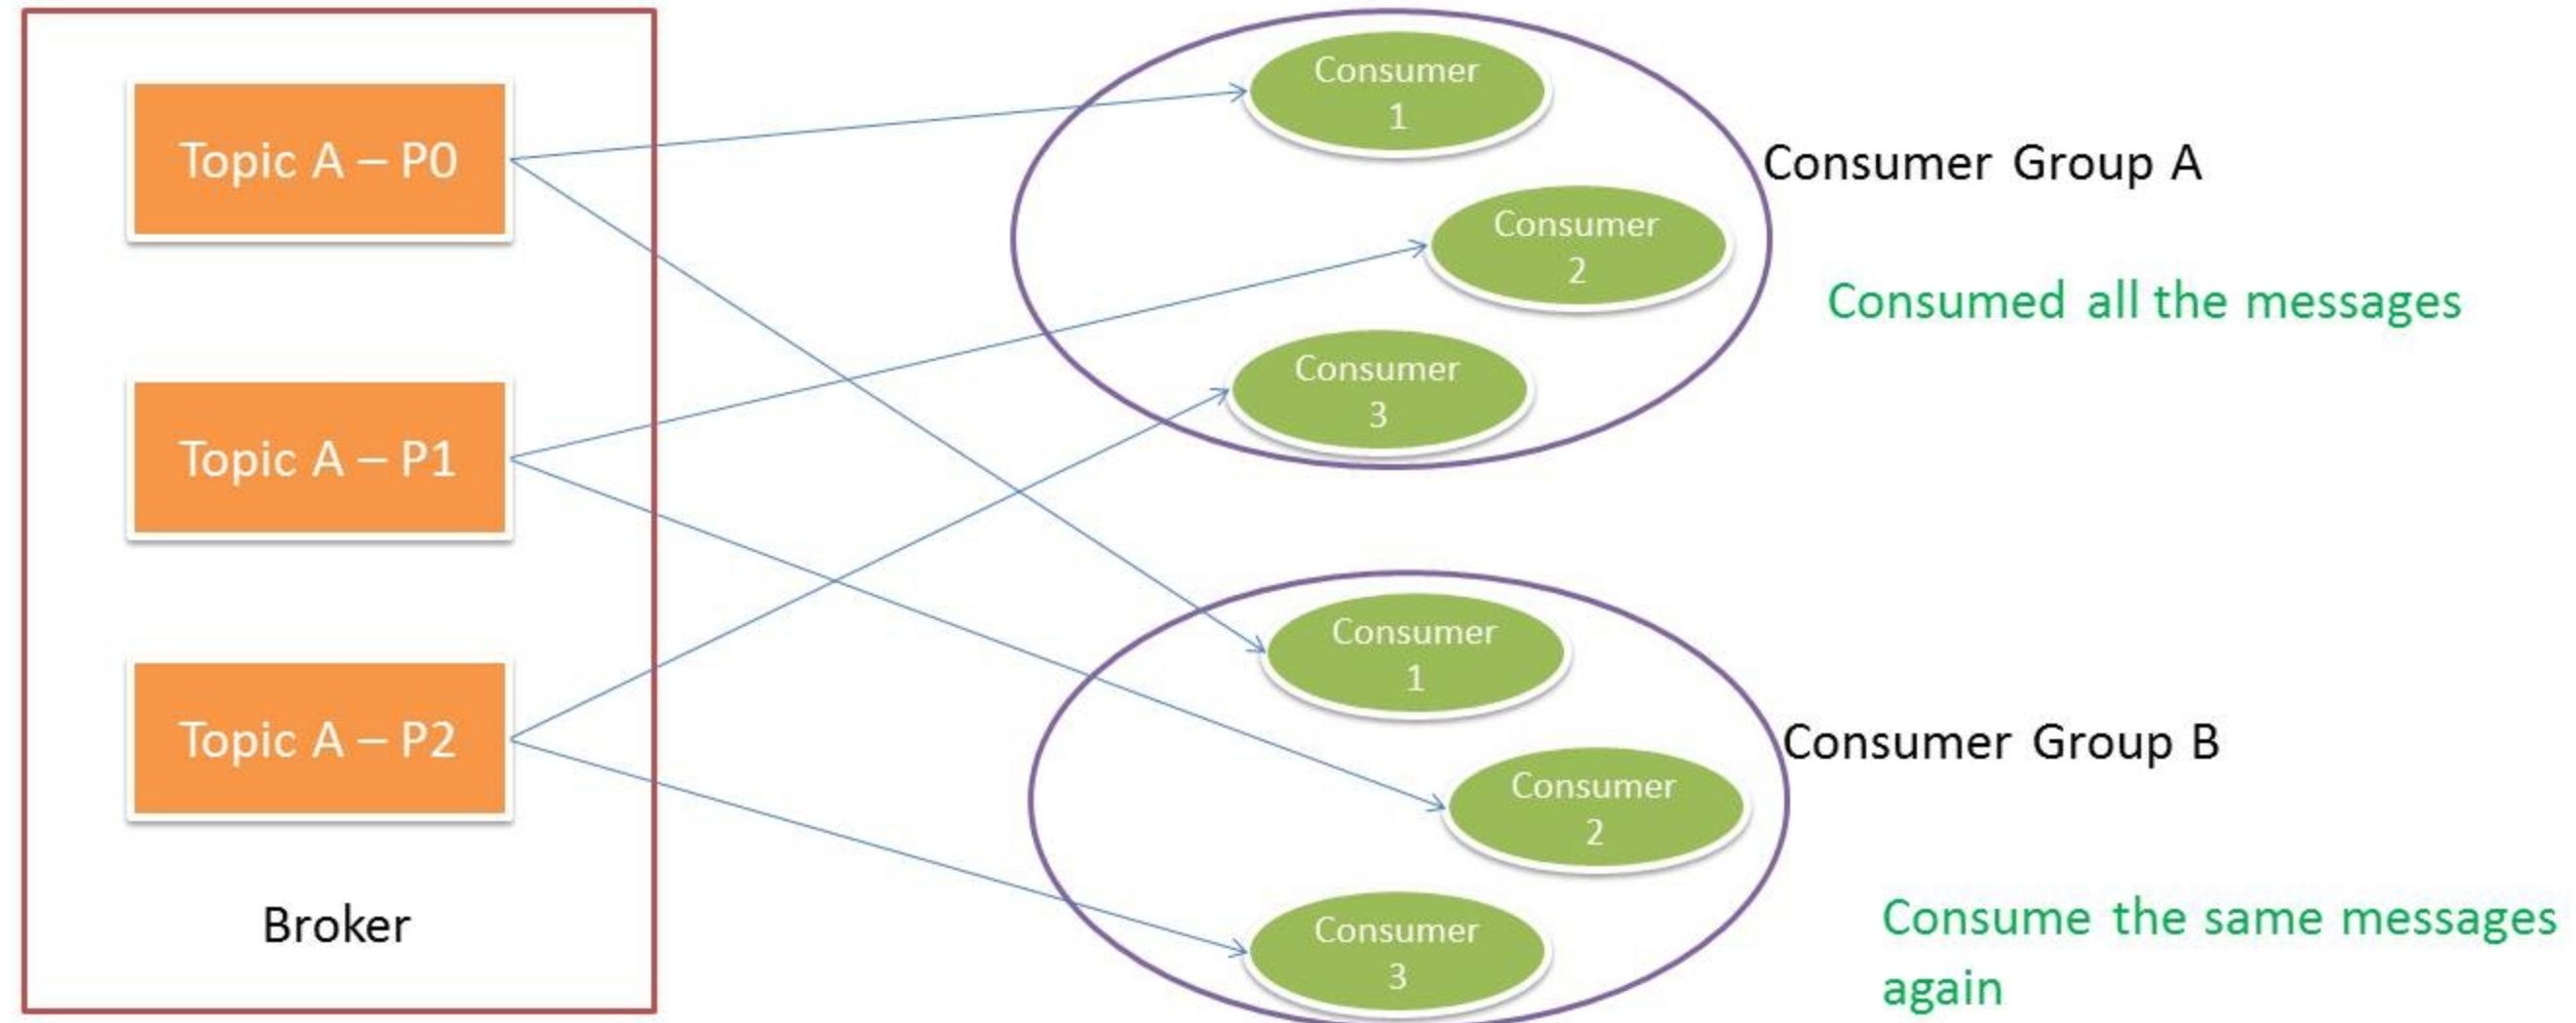
group.id=consumer_group_1

auto.commit.interval.ms=1000

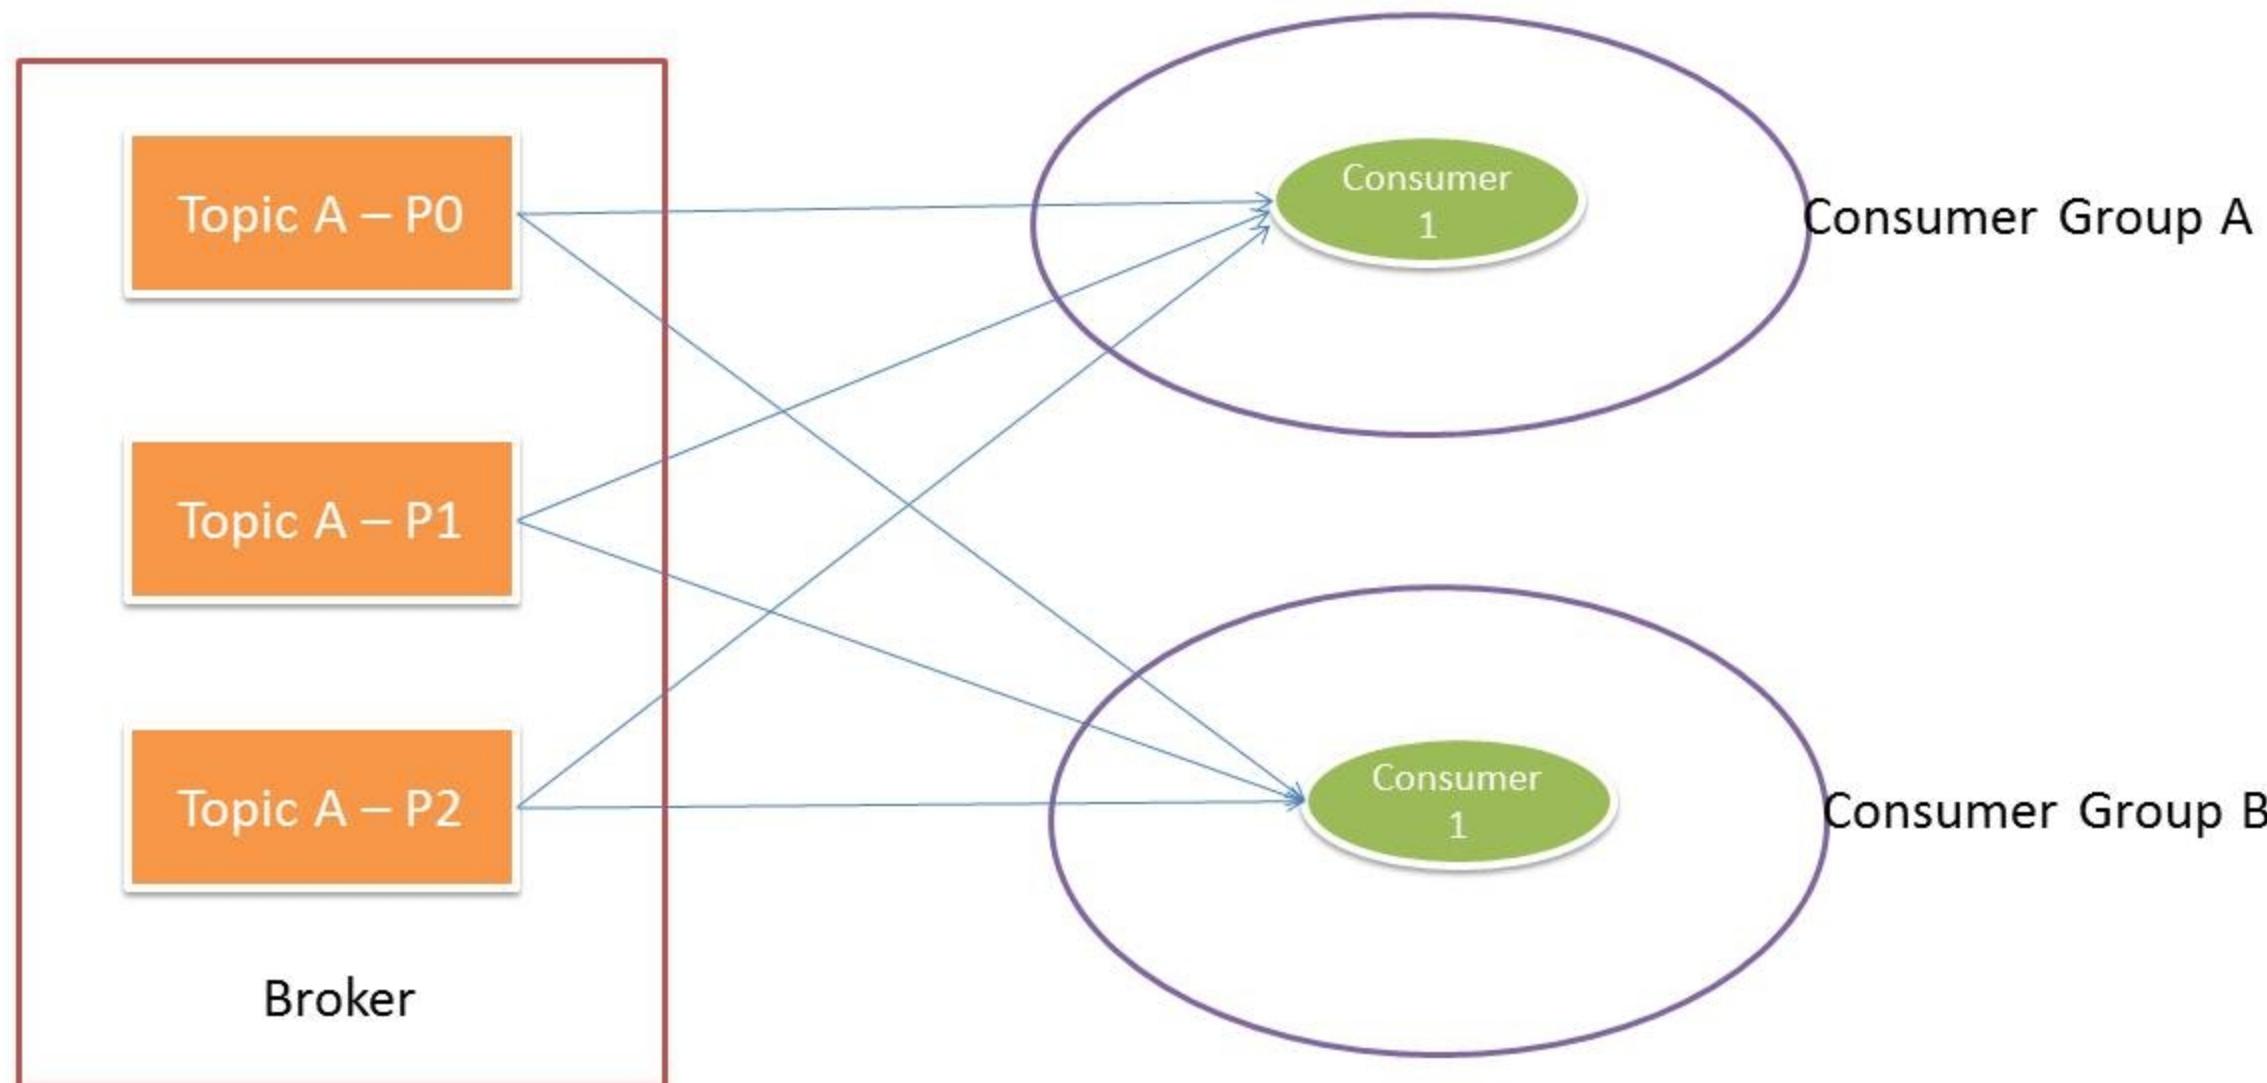
- The property “**zookeeper.connect**” is used to identifies the zookeeper cluster. Kafka consumer uses ZooKeeper to store offsets of messages consumed for a specific topic and partition.

- The property “**group.id**” defines the Consumer group.
- The property “**auto.commit.interval.ms**” defines after how many milliseconds the consumed offsets are written to ZooKeeper.

Consumer Group



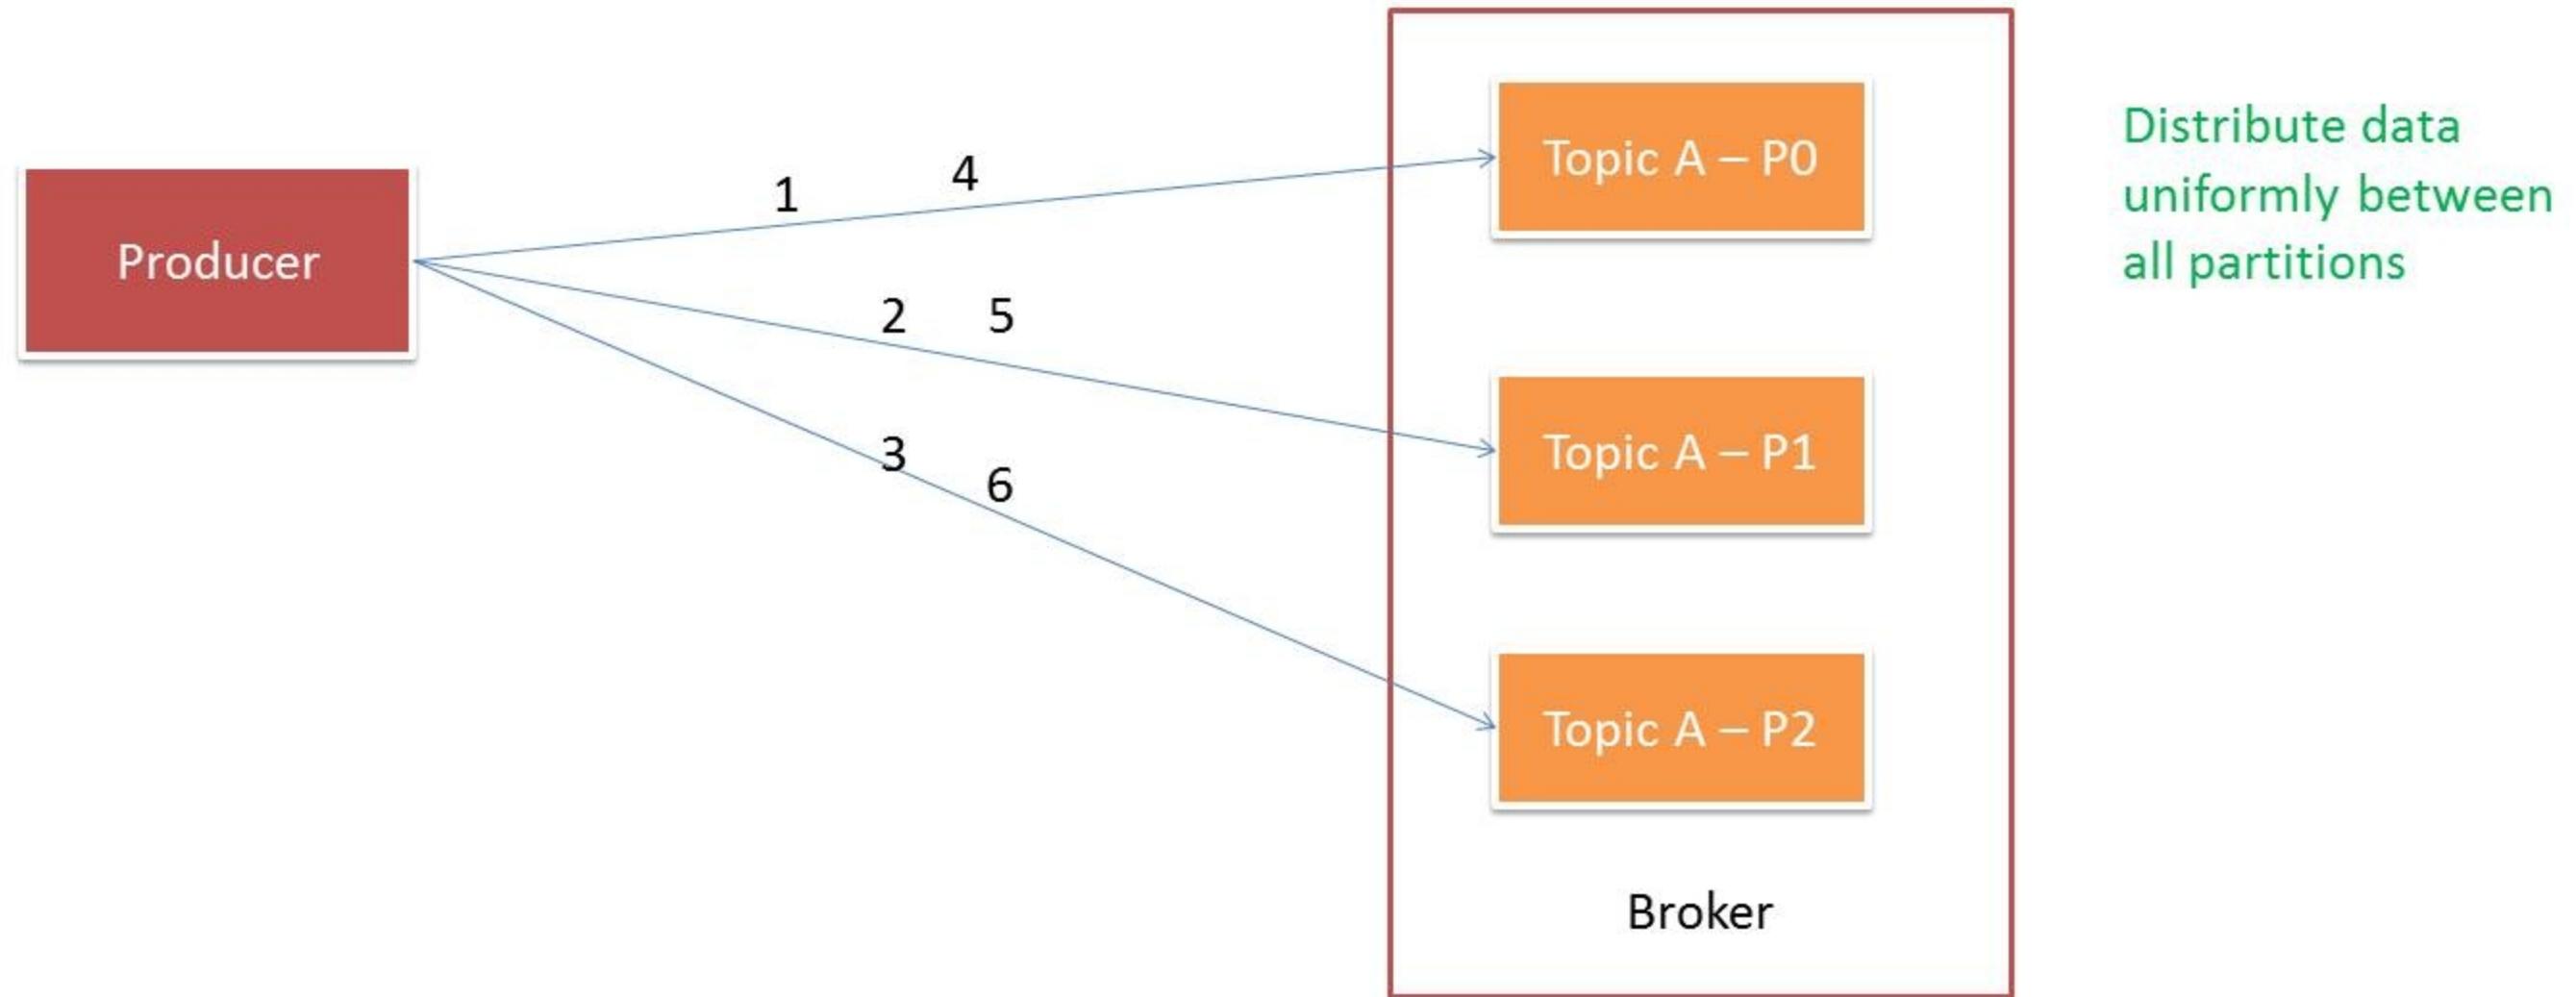
Consumer Broadcast



- Java API to read data from Kafka
- How consumer offset is maintained in zookeeper
- How we can read same data again from kafka

- Need to set the following configuration:
 - `partitioner.class=com.dataflair.CustomPartitioner`
- The property "**partitioner.class**" defines what class to use to determine which Partition in the Topic the message is to be sent to.
- We can define the partition class by implement the **partition(Object key , int paritionCount)** method of **kafka.producer.Partitioner** interface

Producer – RoundRobinPartition



Sample Custom Partition Example – Hands on



- How we can set partition class in kafka.

DataFlair Web Services Pvt Ltd

+91-8451097879

info@data-flair.com

<http://data-flair.com>